University of Basel

Managing and Distributing Software Updates Using Append-Only Logs

Maximilian Barth <m.barth@unibas.com>

Departement of Mathematics and Computer Science

Examiner: Prof. Dr. Christian Tschudin

Supervisor: MSc. Fabrizio Parrillo

12.07.2022

Outline

1. Goal
2. Tinyssb
3. Versioning System
4. Demo
5. Outlook

Scenario

Sensor network

- Sensors are solar powered
- Irregular operating hours
- Nodes communicate via LoRa
- Limited processing power and memory

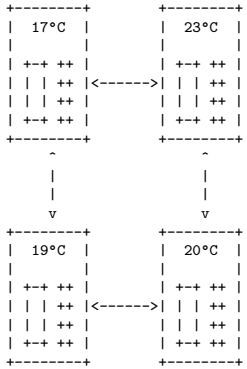

Goal

Distribute and **manage** updates across a solar-powered sensor network.
$\Rightarrow$ consider limitations of LoRa and hardware

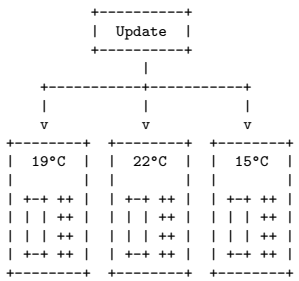

Challenges

Network protocol must allow nodes to **catch up** with missed messages.

- Important for updates
- Still guarantee message authenticity and integrity

$\Rightarrow$ Append-only log protocol **Tinyssb**

github.com/tschudin/tinyssb

Tinyssb

- Append-only logs, so-called **feeds**
- Each feed entry (or packet) is **128B**
- Only the owner of a feed can append new **packets**
- Every node can **verify** a packet's authenticity and integrity
- Nodes request **missing** packets
- Feeds can have **child feeds**

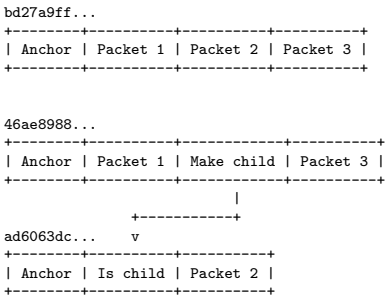

Versioning system

Idea:

- Continuous code deployments
- Allow reversion of updates
- Enable users to create different update branches (similar to Git)
- Provide GUI for interaction with versioning system

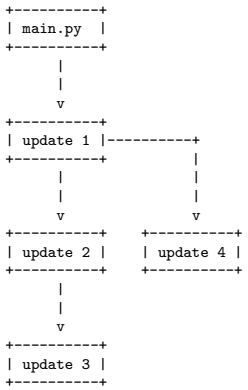

Tinyssb and the versioning system

- The **update feed** contains all the information of the versioning system

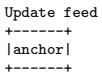

Tinyssb and the versioning system

- The **update feed** contains all the information of the versioning system
- Its first child feed is the **version control feed**, which contains the currently applied version number of each monitored file

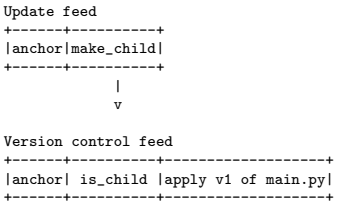

Tinyssb and the versioning system

- The **update feed** contains all the information of the versioning system
- Its first child feed is the **version control feed**, which contains the currently applied version number of each monitored file
- All remaining child feeds are **file update feeds**, each of which contains updates and their dependencies

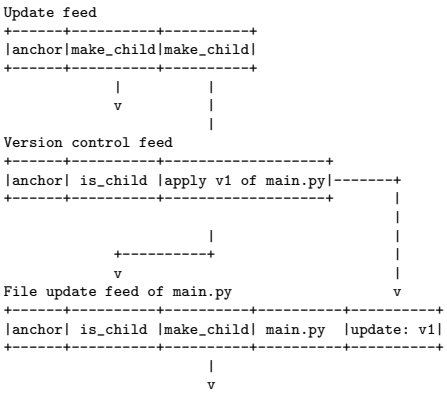

Bad updates

What happens in the case of bad updates?

- Already appended updates cannot be removed
- Large updates may congest the network

⇒ This must be handled

bad_update.c

```
+-------------------------+
|1| int main() {          |
|2|   int x = 0;          |
|3|                       |
|4|   int* ptr = x;       |
|5|   *ptr = 3;           |
|6| }                     |
|7|                       |
+-------------------------+
```

Emergency feeds

Every file has an emergency feed, which can be activated through the following steps:

1. Create a new emergency feed
2. Append the file name to the old emergency feed
3. Append the emergency update

⇒ The emergency feed becomes the new file update feed.

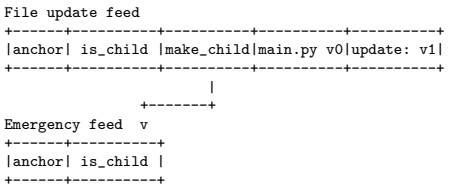

Emergency feeds cont'd

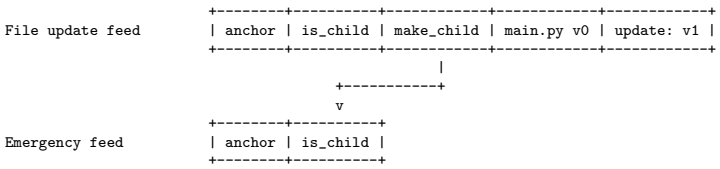

Emergency feeds cont'd

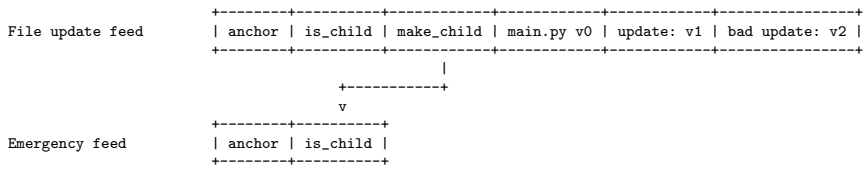

Emergency feeds cont'd

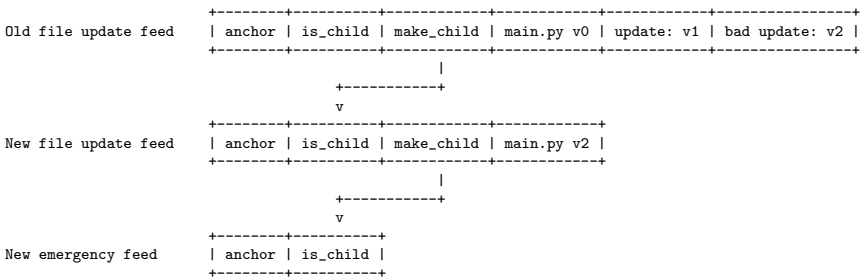

Emergency feeds cont'd

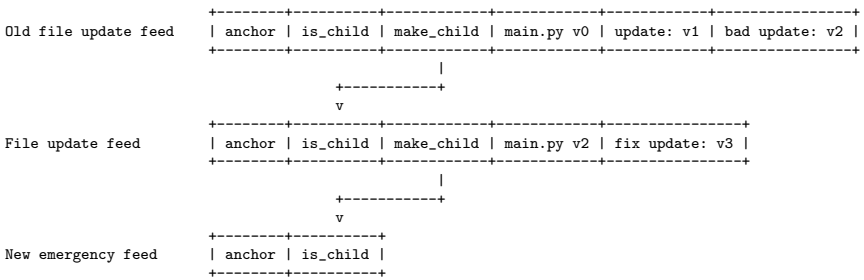

$\Rightarrow$ ignore bad update

Representing updates

How should updates be encoded?
$\Rightarrow$ consider low memory and data rate

Three approaches:

1. Send the entire new file
2. Only send lines that have changed
3. Only send substrings that have changed

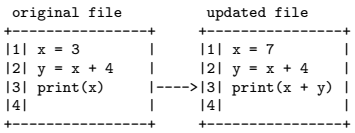

Representing updates cont'd

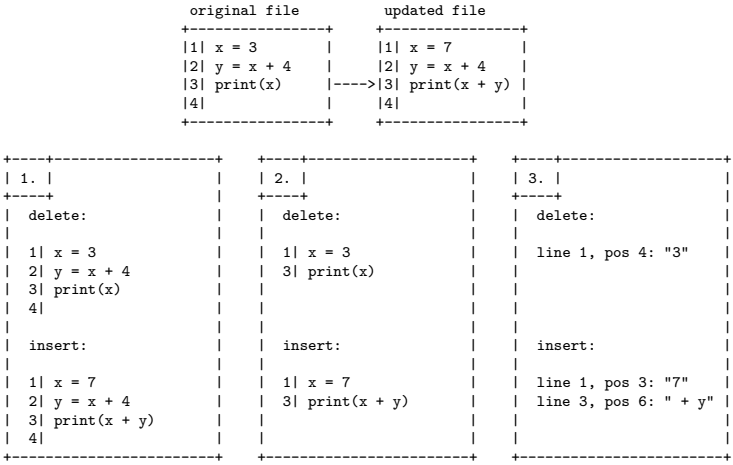

LCS

Idea of using the longest common subsequence **(LCS)** problem to determine the necessary insert and delete operations.

- Same approach as in UNIX's `diff` utility, developed by J. W. Hunt and M. D. McIllroy
- Compare LCS with original and updated file
- Results in compact updates

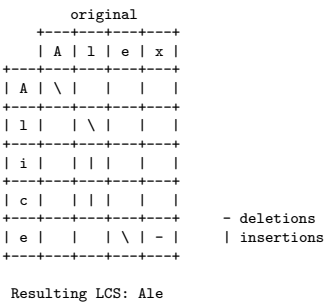

LCS example

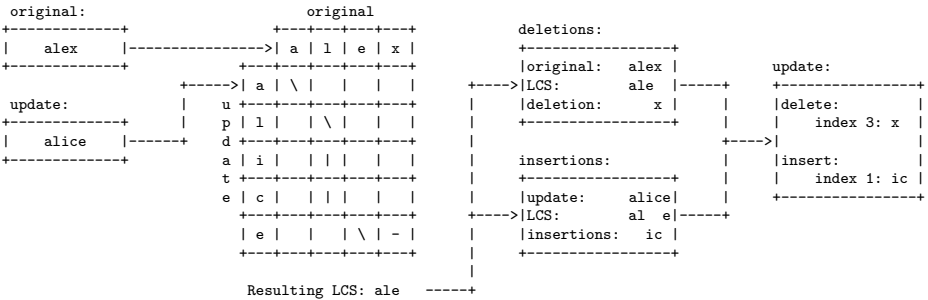
original:
alex
update:
alice
original
a l e x
a \ | | |
l | \ | |
e | | | |
| - | | |
Resulting LCS: ale
deletions:
original: alex
LCS: ale
deletion: x
insertions:
update: alice
LCS: al e
insertions: ic
update:
delete:
index 3: x
insert:
index 1: ic

Managing updates

- The original file is considered version 0
- Each update of a file depends on an already existing version
  $\Rightarrow$ allows creation of different update branches
- Results in a dependency tree

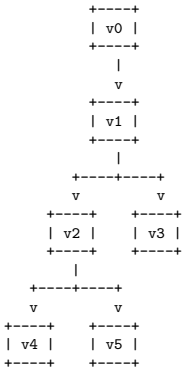

Version jumping

It is possible to jump in between **any** two versions of a file using its dependency tree:

1. **Extract** a path using depth first search
2. **Revert** updates until the latest common predecessor version is reached
3. **Apply** the remaining updates

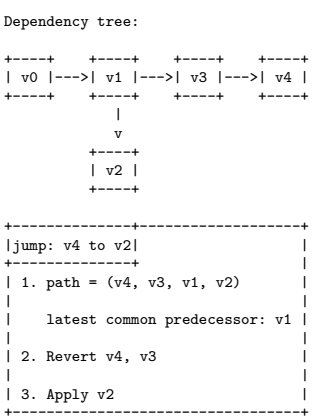

Demo

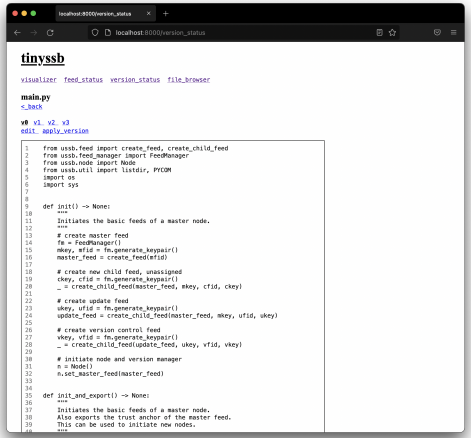
localhost:8000/version_status
localhost:8000/version_status
tinyssb
visualizer feed_status version_status file_browser
main.py
< back
v0 v1 v2 v3
edit apply_version
1 from ssb.feed import create_feed, create_child_feed
2 from ssb.feed_manager import FeedManager
3 from ssb.node import Node
4 from ssb.util import listdir, PYCOM
5 import os
6 import sys
7
8
9 def init() -> None:
10 """
11 Initiates the basic feeds of a master node.
12 """
13 # create master feed
14 fm = FeedManager()
15 mKey, mFid = fm.generate_keypair()
16 master_feed = create_feed(mFid)
17
18 # create new child feed, unassigned
19 ckey, cfid = fm.generate_keypair()
20 _ = create_child_feed(master_feed, mKey, cfid, ckey)
21
22 # create update feed
23 ukey, ufid = fm.generate_keypair()
24 update_feed = create_child_feed(master_feed, mKey, ufid, ukey)
25
26 # create version control feed
27 vkey, vfid = fm.generate_keypair()
28 _ = create_child_feed(update_feed, ukey, vfid, vkey)
29
30 # initiate node and version manager
31 n = Node()
32 n.set_master_feed(master_feed)
33
34
35 def init_and_export() -> None:
36 """
37 Initiates the basic feeds of a master node.
38 Also exports the trust anchor of the master feed.
39 This can be used to initiate new nodes.

Outlook

- Introduce cross-file dependencies
- Integrate with Simon Laube's project
- Field-test in a larger network

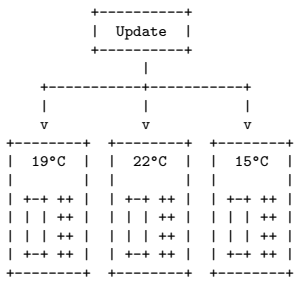

Questions?

m.barth@unibas.ch